

Set 12: Ordered Dictionary Abstract Data Types: (2, 4) Trees, B Trees

CS240: Data Structures and Data Management

Jérémy Barbay

Outline

(2, 4) trees

- Definitions
- Properties
- Insertion
- Deletion

B-Trees

- Definition
- Motivations
- Improvements

Conclusion to Ordered Dictionary ADTs


- Concepts
- References

Multi-way search trees

Definition

A **d-node** is an internal node with

- ▶ d children, t_1, \dots, t_d , and
- ▶ $d - 1$ keys such that $k_1 < k_2 < \dots < k_{d-1}$.




This generalizes binary search trees to larger degrees.

Multi-Way Search Trees

Definition

A **Multi-Way Search Tree** is an ordered search tree consisting of linked d -nodes, where each node may have a different value for d :




It is searched almost like a binary tree.

(2, 4)-Trees

Definition

A **(2,4)-tree** is a multi-way search tree such that

- ▶ Every node has between 2 and 4 children
- ▶ All external nodes have the same depth




(2, 4)-Trees

Definition

A **(2,4)-tree** is a multi-way search tree such that

- ▶ Every node has between 2 and 4 children
- ▶ All external nodes have the same depth



Note: As all external nodes have the same depth, placeholders don't carry much information anymore, and can be omitted.

Properties

Theorem


Consider a (2,4)-tree with n internal **keys**:

1. The number of external placeholders is $|E| = n + 1$.
2. The height h is $\Theta(\log n)$.

Proof: Exercise.

Insertion


- ▶ Find deepest node where the key belongs, and insert.
- ▶ If **overflow**, perform a **node split**:




- ▶ The third element (counting the new element) moves up to parent, possibly causing a new overflow.

Insertion


Example



Insert(S)




Insert(F)




Insertion


Example (cont)



Insert(T)




Overflow in RSTU!



Overflow in GLPT!

Insertion

Example (end)




Theorem


Insertion requires at most $h \in O(\lg n)$ operations.

Deletion

- ▶ Search for the key
- ▶ As in an AVL:
 - ▶ if it has no children, remove it.
 - ▶ if it has children, replace it with in-order predecessor or successor.
- ▶ If too few keys (**underflow**), then
 - ▶ **Transfer** a node from an *immediate sibling* if possible




- ▶ Otherwise, **fuse** with an *immediate sibling* and parent element




Deletion

Example




Delete(A) Simply remove A




Deletion

Example (cont)




Delete(J) Replace J with H




Deletion

Example (cont)




Delete(N) Transfer from RS




Deletion

Example (cont)




Delete(L) Fusion MRP and split




Deletion

Example (cont)




Delete(G) Fusion CE + Underflow + transfer from RU




Deletion

Example (cont)




Delete(Z) immediate

Delete(S) Fusion UX + Underflow + Fusion HR



Deletion

Example (end)



Theorem

Deletion requires at most $h \in O(\lg n)$ operations.

Summary for (2,4)-trees

- ▶ search easy
- ▶ insert may involve several **splittings**.
- ▶ deletion may involve one **transfer**, or several **fusions**.
- ▶ h increases only if root is split.
- ▶ h decreases only if root's sibling's fuse, and root becomes empty.
- ▶ $O(\log n)$ since constant amount of work at each node.

Outline

(2,4) trees

Definitions
Properties
Insertion
Deletion

B-Trees

Definition
Motivations
Improvements

Conclusion to Ordered Dictionary ADTs

Concepts
References

B-Trees

Definition

- ▶ A generalization of (2,4)-trees
- ▶ A B-Tree of order d ($d \geq 3$) is a **multi-way search** such that
 - ▶ every node has $\leq d$ children
 - ▶ every *non-root* node has $\geq \lceil \frac{d}{2} \rceil$ children
 - ▶ all the external nodes have the same depth
- ▶ Often called an (a, b) -tree where $a = \lceil \frac{d}{2} \rceil$ and $b = d$
- ▶ The operations are performed the same as before
 - ▶ For overflow we promote element $\lceil \frac{d+1}{2} \rceil$ (counting the new element)


B-Tree of Order 6

Example


Also known as a (3,6)-tree:

BCDEF

- ▶ Insert(A)



- ▶ Delete (F)



- ▶ Delete (B)

ACDF

Motivations

External Searching

- ▶ What if the dictionary cannot fit in main memory?
- ▶ Need to store data in persistent memory (e.g. on disks)
- ▶ A single access to the data structure takes much longer
 - ▶ RAM Seek – 100 000 memory accesses
 - ▶ Disk Seek – for 1 disk access
- ▶ A disk access brings in a whole page of data
- ▶ Assume we can fit B dictionary elements per page:
How many disk accesses would binary search (or an AVL tree) require?

$$\lg n - \lg B$$

B-Trees

- ▶ **Problem:** One disk access only cuts range of keys in half
- ▶ **Solution:** Use a B-Tree of order d
 - ▶ Choose d such that one d -node fills exactly one disk page
 - ▶ One disk access narrows search a lot more
- ▶ Searching the d -node is still far less expensive than bringing it into memory
- ▶ Running time is proportional to the number of blocks read
 - ▶ Insert and delete designed to reduce the number of d -nodes examined

Performance

- ▶ Suppose $d = 256$
- ▶ Minimum and maximum number of keys *found at each depth*:

Depth	Minimum # Keys	Maximum # Keys
0	1	255
1	254	65,280
2	32,512	16,711,680
3	4,161,536	4,278,190,080
4	532,676,608	1.1×10^{12}

Property

Theorem

The height h of a B-tree of order d is

- ▶ $\Omega(\log_d(n))$
- ▶ $O(\log_{\lceil \frac{d}{2} \rceil}(n))$

Proof.

We know $2^{\lceil \frac{d}{2} \rceil h-1} \leq |E| \leq d^h$ combined with $|E| = n + 1$ gives the result. \square

Improvements

One-Pass Update

One-Pass Update

- ▶ Insert and delete require two passes
 - ▶ First pass down tree finds the bottom level node
 - ▶ Second pass up tree performs splitting or fusing
- ▶ Algorithm can be reworked to perform preemptive splits and fusions on the way down
- ▶ See CLRS textbook for details
 - ▶ You will never need to perform this in our class

Improvements

B*-Tree

B*-Tree

- ▶ Each non-root *d*-node could have as low as 50% utilization
- ▶ On average a node is 69% filled
- ▶ We could insist that a non-root node is at least $\frac{2}{3}$ filled
- ▶ More difficult to do a node split or fusion
- ▶ A B*-Tree node is 90% filled on average

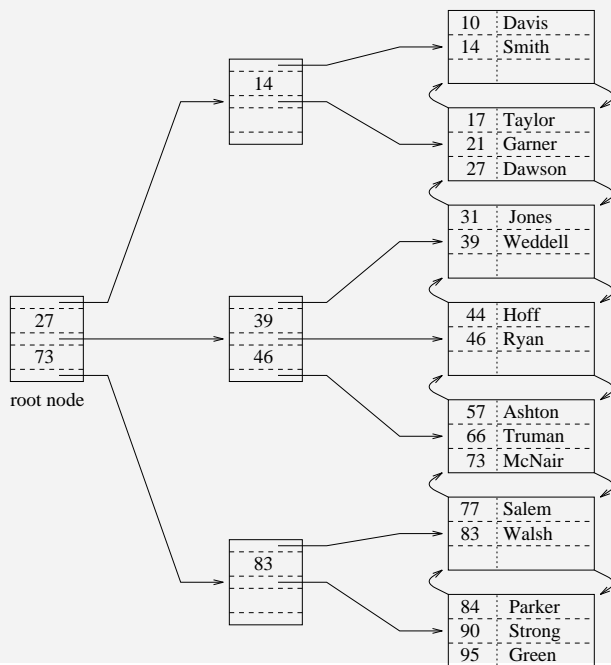
Improvements

B⁺-Tree

B⁺-Tree

- ▶ Desire greatest branching possible
- ▶ Internal nodes contain only keys (not the corresponding satellite data)
- ▶ Bottom level of tree contains the real key-data pair
- ▶ We also wish to perform **Range Queries**
 - ▶ List all professors with id-numbers between 39 and 75
- ▶ Each node has a pointer to the next and previous bottom level page

B⁺-Tree



Summary for B-trees

- ▶ Generalisation of (2,4)-trees
- ▶ Useful for **large** dictionaries, which does not fit in memory.
- ▶ **Several variants**, corresponding to various needs.
- ▶ Very important in practical Databases.

Outline

(2,4) trees

- Definitions
- Properties
- Insertion
- Deletion

B-Trees

- Definition
- Motivations
- Improvements

Conclusion to Ordered Dictionary ADTs

- Concepts
- References

Ordered Dictionary ADTs and their DS

- ▶ Array **Compact**
- ▶ Binary Search Tree (BST) **Fast for static**
- ▶ Sequence (Skip Lists) **Dynamic**
- ▶ AVL **Fast and Dynamic**
- ▶ (2,4) Trees **Fast and “amortized” dynamic**
- ▶ B-Trees **extension for memory issues.**

Diferent solutions to different problems...

References

	GT	CLRS
Arrays	pp. 140-151	pp. 253-264
BST		
Skiplists	pp. 195-202	Not covered.
AVL	pp. 152-158	pp. 296 (poorly covered)
(2,4)-trees	pp. 159-169	pp. 434-452 (indirectly)
B-Trees	pp. 649-653	pp. 434-452