

# Valued Dictionary Abstract Data Types: Databases - SQL

CS240: Data Structures and Data Management  
Slide Set 14

Jérémy Barbay

March 7-9 2006

## Outline

SQL

Implementation

## Outline

- ▶ We will spend about two lectures on databases
  1. Overview of databases
  2. Terminology
  3. SQL
  4. Relational Algebra
  5. Implementation

## Overview

- ▶ A **large** amount of **structured** data stored for a **long** time
  - ▶ Gigabytes and Terabytes of information
  - ▶ Generally outlasts the original technology
- ▶ Updates are frequent
- ▶ Most of the data exists indefinitely
- ▶ Nature of searches change over time
- ▶ Utilize a **DBMS** – Database Management System
- ▶ Only consider **relational databases**

## Peanuts Database

### ► Students Relation:

StudentID	Name	Address	Phone
12345	C. Brown	12 Apple St.	555-1234
67890	L. Van Pelt	34 Pear Ave.	555-5678
22222	P. Patty	56 Grape Blvd.	555-9999

### ► Grades Relation:

Course	StudentId	Grade
CS101	12345	85
CS101	67890	70
EE200	12345	65
EE200	22222	75
CS101	33333	83
PH100	67890	69

## Peanuts Database

### ► Courses Relation:

Course	Prerequisite
CS101	CS100
EE200	EE005
EE200	CS100
CS120	CS101
CS121	CS120
CS205	CS101
CS206	CS121
CS206	CS205

### • Times Relation:

Course	Day	Hour
CS101	M	9AM
CS101	W	9AM
CS101	F	9AM
EE200	T	10AM
EE200	W	1PM
EE200	R	10AM

### ► Rooms Relation:

Course	Room
CS101	Turing Aud.
EE200	25 Ohm Hall
PH100	Newton Lab

## Terminology

- **Attribute** – Name of a set
  - Column name
- **Tuple** – Ordered set; maps attributes to values
  - One row of a table
- **Relation** – Finite set of tuples
  - The table
  - This is the actual information
- **Relation Scheme** – Set of attribute names in a relation
  - Set of column names for a table
- **Database** – Set of relations
  - Set of tables
- **Database Scheme** – Set of relation schemes in a database
  - All the attributes and their various relationships

## Assumptions

- No duplicate rows
- Exactly one entry in any row  $i$  and column  $j$  of a relation
- Rows **are not** ordered
- Columns **are** ordered

## Queries

- ▶ We will focus on lookups/queries primarily
- ▶ More powerful requests than standard Find( key )
  - ▶ **Random Access** – find tuple with key  $K$
  - ▶ **Sequential Access** – a list of all the tuples
  - ▶ **Value** – find tuple with attribute  $A = v$
  - ▶ **Range** – find all tuples with attributes in given range
  - ▶ **Function** – compute function  $f$  for an attribute over a relation
  - ▶ **Boolean** – find all tuples satisfying a boolean expression
  - ▶ **Quantified** – give a yes or no answer
  - ▶ **Similarity** – tuples matching a regular expression

## SQL

- ▶ General form of the SQL SELECT statement:  

```
SELECT
( DISTINCT          —  $\varepsilon$  )
( columns — * — |aggregate function| |column| )
FROM relations
( WHERE conditions —  $\varepsilon$  )
( GROUP BY columns —  $\varepsilon$  )
( ORDER BY columns —  $\varepsilon$  )
```
- ▶ *conditions* consists of attribute names connected by logical and arithmetic operators
- ▶ *aggregate function* can be:
  - ▶ AVG()
  - ▶ MIN()
  - ▶ MAX()
  - ▶ COUNT()
  - ▶ SUM()

## Queries

- ▶ **Note:** The result of a query is a new (unnamed) relation
- ▶ List the personal information for all of the students sorted by student number:

### Solution

```
SELECT Name, Address, Phone FROM Students
ORDER BY StudentID;
```

- ▶ List all student numbers for which a grade has been recorded:

### Solution

```
SELECT StudentID FROM Grades;
```

- ▶ List all unique student numbers for which a grade has been recorded:

### Solution

```
SELECT DISTINCT Student ID FROM Grades;
```

## Queries

- ▶ What is the highest student number?

### Solution

```
SELECT MAX(StudentID) FROM Students;
```

- ▶ Provide a nicer column name for the above query.

### Solution

- ▶ What is Charlie Brown's phone number?

### Solution

```
SELECT Phone FROM Students WHERE Name = 'C.Brown';
```

- ▶ What is the average grade in each course?

### Solution

```
SELECT Course, AVG(Grade) FROM Grades GROUP BY
Course;
```

## Cartesian Product

- ▶ We saw the set operators **union**, **intersection** and **difference**
- ▶ We can also take the **Cartesian product**
- ▶ Every tuple in first set is combined with every tuple in the second set

$R =$	<table><tr><th>A</th><th>B</th></tr><tr><td>0</td><td>1</td></tr><tr><td>2</td><td>3</td></tr></table>	A	B	0	1	2	3	$S =$	<table><tr><th>A</th><th>B</th></tr><tr><td>0</td><td>1</td></tr><tr><td>4</td><td>5</td></tr></table>	A	B	0	1	4	5	$T =$	<table><tr><th>C</th><th>D</th><th>E</th></tr><tr><td>6</td><td>7</td><td>8</td></tr></table>	C	D	E	6	7	8
A	B																						
0	1																						
2	3																						
A	B																						
0	1																						
4	5																						
C	D	E																					
6	7	8																					

▶  $R \times T$ :

$R \times S$ :

## Queries

- ▶ What grade did each student receive (using their name, not their student number)?

### Solution

```
SELECT R.Name, S.Course, S.Grade FROM Students R,  
Grades S  
WHERE R.StudentId=S.StudentId;
```

- ▶ What prerequisites must Charlie Brown have taken?

### Solution

```
SELECT DISTINCT C.Prerequisite  
FROM Students S, Grades G, Courses C  
WHERE S.name = 'C. Brown'  
AND S.StudentID = G.StudentID  
AND G.course=C.course;
```

## Queries

- ▶ Suppose we had the following **CS240Grades Relation**:

StudentId	Grade
98765	85
⋮	⋮

- ▶ What student numbers got above average in the course?

### Solution

```
SELECT StudentId  
FROM CS240Grades C1,  
(SELECT AVG(Grade) AS A  
FROM CS240Grades) C2  
WHERE C1.Grade > C2.A
```

## Summary

- ▶ Databases are one of the **main motivation** for efficient Data Structures.
- ▶ **Sorted data** is helping, but not always available (or is sorted in another order).

## Outline

SQL

Implementation

## Relational Algebra

- ▶ Allows us to build mathematical expressions for relational queries
- ▶ Express query without giving details of how to implement
- ▶ Allows for optimizations
  - ▶ One SQL statement may correspond to numerous expressions
  - ▶ Use algebraic laws to convert expressions
  - ▶ Apply rules to (hopefully) minimize total work

## Select

- ▶  $\sigma_C(R)$
- ▶ Select tuples from relation  $R$  satisfying condition  $C$
- ▶  $C$  can be connected by arithmetic and boolean operators
- ▶ Horizontal subset
- ▶ Like the `WHERE` clause in SQL
- ▶ **Example:** List all the tuples for grades reported in course CS101

Solution

$\sigma_{Course="CS101"}(Grades)$

## Project

- ▶  $\pi_A(R)$
- ▶ Delete from relation  $R$  all attributes not in  $A$
- ▶ Implicitly removes duplicate rows
- ▶ Vertical subset
- ▶ Like the `SELECT DISTINCT` clause in SQL
- ▶ **Example:** List all the students and their phone numbers

Solution

$\pi_{Name,Phone}(Students)$

- ▶ We can perform composition between various operators
- ▶ **Example:** List all student numbers having taken CS101

Solution

$\pi_{StudentID}(\sigma_{Course="CS101"}(Grades))$

## Join

- ▶  $R \bowtie_a S$
- ▶ Concatenate all tuples from  $R$  and  $S$  where  $R.a = S.a$
- ▶ Removes duplicate columns
- ▶ Natural Join –  $R \bowtie S$ 
  - ▶  $R$  and  $S$  have only one column in common
- ▶ Combines FROM and WHERE clause in SQL
- ▶ **Example:** List the time and room for each class

### Solution

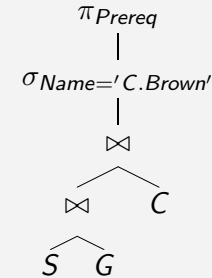
$Rooms \bowtie Times$

## Representation

- ▶ What prerequisites must Charlie Brown have taken?
- ▶ Directly translate the SQL from last lecture:

$$\pi_{Prereq} \left( \sigma_{Name='C.Brown'} ((S \bowtie G) \bowtie C) \right)$$

- ▶ We can draw an expression tree:




## An Optimization

- ▶ An equivalent representation:

$$\pi_{Prereq} \left( \pi_{Course} \left( \sigma_{Name='C.Brown'} (S) \right) \bowtie \pi_{\{Course, Id\}} (G) \right) \bowtie C$$

- ▶ Corresponding expression tree:



## How to Optimize

### 1. Optimize the expression tree

- ▶ Usually,  $\text{Cost}(\bowtie) \geq \text{Cost}(\pi) \geq \text{Cost}(\sigma)$
- ▶ We want to push selection and projection as far down the tree as possible
- ▶ Apply laws such as:
  - ▶  $\sigma_c(R \bowtie S) \equiv \sigma_c(R) \bowtie S$  if  $c \notin S$
  - ▶  $\pi_A(R \bowtie S) \equiv \pi_A(\pi_{A'}(R) \bowtie \pi_{A''}(S))$
- ▶ Duplicate detection should be left until the end
- ▶ Use heuristics to decide when to apply rules

### 2. Optimize the implementation of each operator

- ▶ Our focus

## Database Implementation

- ▶ Still treat a database as a dictionary ADT
- ▶ **Primary Key** – One or more fields whose value(s) uniquely identify each tuple in a relation
- ▶ **Primary Index** – Main data structure organized by primary keys
  - ▶ B<sup>+</sup>-Tree
  - ▶ Extendible Hashing
  - ▶ What is 12345's phone number?
- ▶ **Secondary Index** – An index based on another attribute
  - ▶ What is Charlie Brown's phone number?
- ▶ Updates are more expensive, but certain queries are much faster

## Indices

- ▶ The secondary index does not store data
- ▶ Points to where the data is in the primary index
- ▶ Example:
  - ▶ Id is the primary key with the index stored in a hash table
  - ▶ Name is the secondary index stored in a B-Tree

## Implementing Operators

- ▶ Generally either sorting based or hashing based
- ▶ Choice depends on size of relation and existing indices
- ▶ **Union**
  - ▶ Sort both relations
  - ▶ Merge removing duplicates
  - ▶ What if you know there are no duplicates?
- ▶ **Difference**
  - ▶ Identical

## Simple Operators

- ▶ **Intersection**
  - ▶ Still much the same
  - ▶ Usually implemented as special case of the join operator
- ▶ **Projection**
  - ▶ Scan each tuple
  - ▶ Removing duplicates is primary problem
  - ▶ Sort as you perform the scan
  - ▶ Could also utilize a hash table

## Selection

- ▶ Duplicate detection not an issue
- ▶ **Equality Selection**
  - ▶  $\sigma_{A=v}$
  - ▶ Use (extendible) hashing
  - ▶ Consider Students with Id as primary index
    - Good:**  $Id = 12345$
    - Bad:**  $Id > 12345$
    - Bad:** Name = 'C.Brown'

## Selection

- ▶ **Range Selection**
  - ▶  $\sigma_{v \leq A \leq w}$
  - ▶ Use B<sup>+</sup>-Tree
  - ▶ Consider Grades with (Course, Id) as primary index
    - Good:** Course = CS101 & Id = 12345
    - Good:** Course = CS101
    - Bad:** Id = 12345
- ▶ Chose the appropriate index based on particular request
- ▶ Create indices based on the queries users tend to ask

## Join

- ▶ Suppose we have the following relations (with sample data):

▶ **R:**

A	B
0	6
2	1
4	3

• **S:**

B	C
6	4
3	8
9	12
6	15

- ▶ Perform  $R \bowtie S$

A	B	C
0	6	4
0	6	15
4	3	8

- ▶ Let  $r = |R|$  and  $s = |S|$

## Join

There are several ways to implement Joins:

- ▶ **Nested Loop**
- ▶ **Sorted Join**
- ▶ **Index Join**

What are the advantages of each implementation?



## Join

**nested loop:** very costly ( $\text{Cost} = rs$ )

```
for each tuple a in R
  for each tuple b in S
    if a.B = b.B then output a, b
```

## Join

**sort join:** merge, mark as  $R$  or  $S$  and sort on  $B$

merge  $R$  and  $S$   
sort on  $B$  attribute  
go through linearly, joining on grouped elements (since they will be sorted together)

$(\text{Cost} = (r + s) \log(r + s))$

## Join

**index join:** if index on join element  $B$  ( $\text{Cost} = r + s$  if index is  $O(1)$  time to find tuples (hash table).  $\text{Cost} = r + s \log s$  if B-tree used on index)

```
for each tuple a in R
  use index of S to find tuples b where b.B=a.B
  for each tuple b in S
    output a, b
```

## Summary

- ▶ Database is a **whole field of research**, we barely scratched the surface here.
- ▶ Database is one of the main motivations for efficient data-structures: how to manage **very large** amount of data.
- ▶ How to compute things is not the whole story, but also in **which order** to compute it.