

Slide Set 17: Memory/Heap Management

CS240: Data Structures and Data Management

Jérémy Barbay

1/ 39

Outline

Memory Management

Principle of Locality

Block Replacement Strategies

Heap Management

Fixed Sized Blocks

Variable Sized Blocks

Buddy Block

Final Review

2/ 39

Goals

- ▶ Minimize expensive memory accesses (disks)
 - ▶ We saw data structures and algorithms that address this
B-Trees
 - ▶ What can the OS do to help?
- ▶ Support dynamic memory allocation
 - ▶ Manage the program **heap**
 - ▶ Treat the heap as an ADT

3/ 39

Principle of Locality

- ▶ Keep important things as near to the CPU as possible
- ▶ A program does not use all pieces of disk/memory equally
- ▶ **Spatial Locality**
 - ▶ Access something at position p
 - ▶ More likely to access something at position $p + \epsilon$
- ▶ **Temporal Locality**
 - ▶ Access something at time t
 - ▶ More likely to access it again at time $t + \epsilon$

4/ 39

Exploiting the Tendencies

- ▶ Break external memory into blocks
- ▶ Spatial Locality
 - ▶ Ask for any address in a block and the whole block is brought in
- ▶ Temporal Locality
 - ▶ **Virtual Memory**
 - ▶ Keep a directory of all external blocks
 - ▶ Bring a block into memory only when accessed
 - ▶ Flag which blocks are in disk cache

5/ 39

Block Replacement

- ▶ What if the disk cache fills?
- ▶ Program accesses memory in block B
Access(B)
 - if** B in disk cache **then**
 Perform access
 - else if** there exists free block F in cache **then**
 Fetch B into F
 Perform access
 - else**
 Evict a block E from cache
 Fetch B into E
 Perform access
 - end if**

6/ 39

Example

- ▶ Suppose 3 internal blocks in cache and 8 external blocks
- ▶ Keep a directory of internal and external blocks
- ▶ Initially empty cache
- ▶ Access blocks 3, 6, 3, 1

External								Internal		
0	1	2	3	4	5	6	7	0	1	2
–	–	–	–	–	–	–	–	–	–	–
–	–	–	0	–	–	–	–	3	–	–
–	–	–	0	–	–	1	–	3	6	–
–	–	–	0	–	–	1	–	3	6	–
–	2	–	0	–	–	1	–	3	6	1

7/ 39

Eviction Policies

- ▶ Suppose we now access block 0
- ▶ How do we determine which block to evict?
- ▶ Goal is to minimize total number of disk fetches
- ▶ *MA* (Memory Access) – Time to perform a memory access to block B
- ▶ *BR* (Block Replacement) – Time to determine block to evict and perform replacement

8/ 39

Random

- ▶ Randomly select an internal block to evict
- ▶ Example – 1, 2, 5, 4, 5, 3, 2, 3 (generator gives 1, 0, 0, 2, 1)

External								Internal		
0	1	2	3	4	5	6	7	0	1	2
–	0	1	–	–	2	–	–	1	2	5

- ▶ $MA - O(1)$
- ▶ $BR - O(1)$
- ▶ Space – 0
- ▶ Worst case scenario – Request the element you just removed.

9/ 39

FIFO

- ▶ Evict block that has been there the longest
- ▶ Example – 1, 2, 5, 4, 1, 5, 2, 5

External								Internal		
0	1	2	3	4	5	6	7	0	1	2
–	0	1	–	–	2	–	–	1	2	5

- ▶ $MA - O(1)$
- ▶ $BR - O(1)$ (with a cyclic order)
- ▶ Space – $O(1)$ (to remember where you are in the cycle)
- ▶ Worst case scenario – $(1, 2, 3, 4)^*$ if $Q = 3$

10/ 39

LFU

- ▶ Evict the least frequently used block in the cache
- ▶ Need to keep a count with each internal item
 - ▶ Keep a PQ of the counts
- ▶ If you access an internal item, update count
- ▶ $MA - O(1)$ (with a trick to normalize frequencies)
- ▶ $BR - O(\lg m)$
- ▶ Space – $O(n)$ (overhead for counting)
- ▶ Worst case scenario – $(1, 2, 3, 4)^*$
- ▶ Comments – Difference between FIFO and LFU?
LFU (but not FIFO) updates its data counter when a block in the cache is queried.

11/ 39

Example

- ▶ Access blocks 4, 1, 5, 5

External								Internal		
0	1	2	3	4	5	6	7	0	1	2
–	1	–	–	0	2	–	–	4	1	5

- ▶ Access block 4

0	1	2	3	4	5	6	7	0	1	2
–	1	–	–	0	2	–	–	4	1	5

- ▶ Access block 2

0	1	2	3	4	5	6	7	0	1	2

12/ 39

LRU

- ▶ Evict the least **recently** used block in the cache
- ▶ Whatever item was accessed furthest in the past
- ▶ Need to keep a queue of internal items
 - ▶ Front of queue is the least recently accessed
- ▶ If you access an internal item, move it to back of queue
- ▶ $MA - O(1)$
- ▶ $BR - O(n)$
- ▶ Space – $O(m)$ or $O(n)$, depends of the variant.
- ▶ Worst case scenario – $(1, 2, 3, 4)^*$ for $Q = 3$.
- ▶ Comments – This is used a lot at the lowest level.

13/ 39

Example

- ▶ Access blocks 4, 1, 5, 1, 4, 2

External								Internal		
0	1	2	3	4	5	6	7	0	1	2
–	1	–	–	0	2	–	–	4	1	5

14/ 39

Summary

- ▶ Various ways to manage Memory.
- ▶ It is always a trade-off between speed and efficiency.
- ▶ It has a **big impact** on the real performance.

Reference:

- ▶ Goodrich and Tamassia, 646–648 and 668–669
- ▶ Aho and Ullman, 146–159
- ▶ most OS books.

15/ 39

Outline

Memory Management

Principle of Locality
Block Replacement Strategies

Heap Management

Fixed Sized Blocks
Variable Sized Blocks
Buddy Block

Final Review

16/ 39

Operators

- ▶ Treat the heap as an ADT
- ▶ Two operators:
 - ▶ $p \leftarrow \text{Allocate}(\text{size})$
 - ▶ $\text{Free}(p)$

17/ 39

Issues

- ▶ Fixed or variable sized blocks?
- ▶ If fixed, small or large sized blocks?
- ▶ User or OS freeing memory blocks?
Explicit vs Implicit freeing of memory.
- ▶ Time or memory?
- ▶ Can blocks reference each other?

18/ 39

Fixed Sized Blocks

- ▶ Treat memory as an array of blocks
- ▶ Some are in use, some are on a free list
- ▶ User can **only** ask for one block at a time
- ▶ Allocation – grab any block from the free list
- ▶ Free – when no longer needed, return to free list

19/ 39

Fixed Sized Blocks

Issues

- ▶ **Internal Fragmentation**
If user does not need all space in block, the extra space is wasted
- ▶ How do we know when block is no longer needed?
- ▶ **Explicit Free**
Simplest to implement

20/ 39

Fixed Sized Blocks

Implicit Freeing

- ▶ **Implicit Free** with reference counts
 - ▶ Update number of references to an allocated block
 - ▶ Islands of garbage may result
- ▶ **Implicit Free** with garbage collection
 - ▶ Periodically clean memory
 - ▶ 2-pass Mark and Sweep algorithm

21/ 39

Variable Sized Blocks

- ▶ Allocate a contiguous chunk of memory
- ▶ All the issues with fixed sized blocks and more
- ▶ **External Fragmentation**
 - ▶ Division of free space leaves no chunk large enough
 - ▶ May require expensive **Memory Compaction**
- ▶ **Internal Fragmentation**
 - ▶ To simplify allocation, OS may give a larger block than was requested
 - ▶ For example, round up to a multiple of a minimum block size

22/ 39

Variable Sized Blocks

Implementation

- ▶ Maintain an ordered list of each chunk of memory
- ▶ **Allocate:**
 - ▶ Search for a free hole that is “big enough”
 - ▶ Mark node as allocated, updating neighbour
- ▶ **Free:**
 - ▶ Mark node as free
 - ▶ Merge with any free neighbours

23/ 39

Variable Sized Blocks


Selection Policies

- ▶ Consider the following sequence on a RAM of size 100:

```
A <- Allocate( 40 )  
B <- Allocate( 40 )  
Free( A )  
C <- Allocate( 10 )
```

- ▶ Various strategies for choosing free block

- ▶ **First Fit**



- ▶ **Best Fit**



24/ 39

Variable Sized Blocks


Comments

- ▶ All strategies can outperform the others in certain cases
- ▶ First-fit and Alternating are simplest to implement
- ▶ Worst-fit needs a priority queue
- ▶ Best-fit needs a more complicated priority queue

25/ 39

Buddy Block

- ▶ Trade less external fragmentation for internal fragmentation
- ▶ OS gives block sizes that are powers of two
 - ▶ Suppose heap has size $N = 2^m$
 - ▶ Valid allocation sizes are $2^0, 2^1, \dots, 2^m$ blocks
 - ▶ Maintain $m + 1$ free lists
 - ▶ Initially all but the 2^m list is empty




26/ 39

Buddy Block

Implementation

- ▶ Only allowed to merge free blocks that are “buddies”
- ▶ Similar to extendible hashing
- ▶ Each block of size 2^k starts at some multiple of 2^k
- ▶ Each block of size 2^k has a buddy of the same size
- ▶ The two buddies are within the same block of size 2^{k+1}



27/ 39

Buddy Block

Operators


- ▶ $p \leftarrow \text{Allocate}(n)$
 - ▶ Round n up to the nearest power of two: 2^i
 - ▶ If no free block of size 2^i , split a block of size 2^{i+1}
 - ▶ This split may require another split ...
- ▶ $\text{Free}(p)$
 - ▶ Add p 's block to the free list of size 2^i
 - ▶ If buddy is in the free list, merge and add to free list of size 2^{i+1}
 - ▶ Again, this may trigger another merge ...

28/ 39


Buddy Block

Example


- ▶ `A <- Allocate(3)`




- ▶ `B <- Allocate(3)`



- ▶ `C <- Allocate(1)`



- ▶ `D <- Allocate(2)`




29/ 39


Buddy Block

Example


- ▶ `Free(B)`



- ▶ `Free(C)`



- ▶ `Free(D)`



30/ 39

Buddy Block

Comments

- ▶ $\Theta(1)$ to find buddy
 - ▶ If free list is implemented with an array
- ▶ $\Theta(m)$ worst case for both Free and Allocate
- ▶ Internal fragmentation wastes approximately 30%
- ▶ External fragmentation still possible

31/ 39

Summary

- ▶ Two extremes and one intermediate solution:
Fixed, Variable Sized Blocks, and **Buddy** Blocks.
- ▶ As for memory management, big impact on performance.

Reference: none in GT nor CLRS.

32/ 39

Outline

Memory Management

- Principle of Locality
- Block Replacement Strategies

Heap Management

- Fixed Sized Blocks
- Variable Sized Blocks
- Buddy Block

Final Review

33/ 39

Final Review

The final can/should cover the following concepts:

- ▶ Algorithmic Analysis
- ▶ ADTs
 - ▶ LIFO, FIFO
 - ▶ Graphs
 - ▶ Trees, ordinal and binary cardinal
 - ▶ Priority Queues and Heaps
 - ▶ Unsorted Dictionaries
 - ▶ Sorted Dictionaries
 - ▶ Valued Dictionaries
- ▶ Applications
 - ▶ Relational Databases and SQL
 - ▶ Compression Algorithms
 - ▶ Pattern Matching Algorithms
 - ▶ Heap and Memory Management.

34/ 39

1. Algorithmic Analysis

- ▶ Asymptotics
- ▶ Sorting algorithms
- ▶ Recursivity (and its analysis)
- ▶ Worst and Average Case complexity
- ▶ Randomized algorithms
- ▶ Lower bounds in the comparison model

35/ 39

2. ADTs

- ▶ LIFO, FIFO
- ▶ Graphs
- ▶ Trees, ordinal and binary cardinal
- ▶ Priority Queues and Heaps
- ▶ Unsorted Dictionaries
 - ▶ with Array
 - ▶ with Lists
- ▶ Sorted Dictionaries
 - ▶ Array
 - ▶ Lists
 - ▶ BST
 - ▶ AVL
 - ▶ (2,4) Trees
 - ▶ B Trees
- ▶ Valued Dictionaries
 - ▶ Hash Table

36/ 39

3. Applications

- ▶ Relational Databases and SQL
- ▶ Compression Algorithms
- ▶ Pattern Matching Algorithms
- ▶ Heap and Memory Management.

37/ 39

Advices

- ▶ Make sure that you know how to do the **assignments**, and the **midterm**.
- ▶ Do the assignments from previous years.
- ▶ Try to remember **Why** and **How** we do things.

38/ 39

Good Luck...

39/ 39