

Slide Set 17: Memory/Heap Management

CS240: Data Structures and Data Management

Jérémy Barbay

Outline

Memory Management

- Principle of Locality

- Block Replacement Strategies

Heap Management

- Fixed Sized Blocks

- Variable Sized Blocks

- Buddy Block

Final Review

Goals

- ▶ Minimize expensive memory accesses (disks)
 - ▶ We saw data structures and algorithms that address this
B-Trees
 - ▶ What can the OS do to help?
- ▶ Support dynamic memory allocation
 - ▶ Manage the program **heap**
 - ▶ Treat the heap as an ADT

Principle of Locality

- ▶ Keep important things as near to the CPU as possible
- ▶ A program does not use all pieces of disk/memory equally
- ▶ **Spatial Locality**
 - ▶ Access something at position p
 - ▶ More likely to access something at position $p + \epsilon$
- ▶ **Temporal Locality**
 - ▶ Access something at time t
 - ▶ More likely to access it again at time $t + \epsilon$

Exploiting the Tendencies

- ▶ Break external memory into blocks
- ▶ Spatial Locality
 - ▶ Ask for any address in a block and the whole block is brought in
- ▶ Temporal Locality
 - ▶ **Virtual Memory**
 - ▶ Keep a directory of all external blocks
 - ▶ Bring a block into memory only when accessed
 - ▶ Flag which blocks are in disk cache

Block Replacement

- ▶ What if the disk cache fills?
- ▶ Program accesses memory in block B

Access(B)

if B in disk cache **then**

Perform access

else if there exists free block F in cache **then**

Fetch B into F

Perform access

else

Evict a block E from cache

Fetch B into E

Perform access

end if

Example

- ▶ Suppose 3 internal blocks in cache and 8 external blocks
- ▶ Keep a directory for internal and external blocks
- ▶ Initially empty cache
- ▶ Access blocks 3, 6, 3, 1

External							
0	1	2	3	4	5	6	7
-	-	-	-	-	-	-	-
-	-	-	0	-	-	-	-
-	-	-	0	-	-	1	-
-	-	-	0	-	-	1	-
-	2	-	0	-	-	1	-

Internal		
0	1	2
-	-	-
3	-	-
3	6	-
3	6	-
3	6	1

Eviction Policies

- ▶ Suppose we now access block 0
- ▶ How do we determine which block to evict?
- ▶ Goal is to minimize total number of disk fetches
- ▶ *MA* (Memory Access) – Time to perform a memory access to block *B*
- ▶ *BR* (Block Replacement) – Time to determine block to evict and perform replacement

Random

- ▶ Randomly select an internal block to evict
- ▶ Example – 1, 2, 5, 4, 5, 3, 2, 3 (generator gives 1, 0, 0, 2, 1)

External							
0	1	2	3	4	5	6	7
–	0	1	–	–	2	–	–

Internal		
0	1	2
1	2	5

- ▶ $MA - O(1)$
- ▶ $BR - O(1)$
- ▶ Space – 0
- ▶ Worst case scenario – Request the element you just removed.

FIFO

- ▶ Evict block that has been there the longest
- ▶ Example – 1, 2, 5, 4, 1, 5, 2, 5

External							
0	1	2	3	4	5	6	7
–	0	1	–	–	2	–	–

Internal		
0	1	2
1	2	5

- ▶ $MA - O(1)$
- ▶ $BR - O(1)$ (with a cyclic order)
- ▶ Space – $O(1)$ (to remember where you are in the cycle)
- ▶ Worst case scenario – $(1, 2, 3, 4)^*$ if $Q = 3$

LFU

- ▶ Evict the least **frequently** used block in the cache
- ▶ Need to keep a count with each internal item
 - ▶ Keep a PQ of the counts
- ▶ If you access an internal item, update count
- ▶ $MA - O(1)$ (with a trick to normalize frequencies)
- ▶ $BR - O(\lg m)$
- ▶ Space – $O(n)$ (overhead for counting)
- ▶ Worst case scenario – $(1, 2, 3, 4)^*$
- ▶ Comments – Difference between FIFO and LFU?
LFU (but not FIFO) updates its data counter when a block in the cache is queried.

Example

- ▶ Access blocks 4, 1, 5, 5

External

0	1	2	3	4	5	6	7
-	1	-	-	0	2	-	-

- ▶ Access block 4

-	1	-	-	0	2	-	-
---	---	---	---	---	---	---	---

- ▶ Access block 2

--	--	--	--	--	--	--	--

Internal

0	1	2
4	1	5

4	1	5
---	---	---

--	--	--

LRU

- ▶ Evict the least **recently** used block in the cache
- ▶ Whatever item was accessed furthest in the past
- ▶ Need to keep a queue of internal items
 - ▶ Front of queue is the least recently accessed
- ▶ If you access an internal item, move it to back of queue
- ▶ $MA - O(1)$
- ▶ $BR - O(n)$
- ▶ Space – $O(m)$ or $O(n)$, depends of the variant.
- ▶ Worst case scenario – $(1, 2, 3, 4)^*$ for $Q = 3$.
- ▶ Comments – **This is used a lot at the lowest level.**

Example

- ▶ Access blocks 4, 1, 5, 1, 4, 2

External

0	1	2	3	4	5	6	7
–	1	–	–	0	2	–	–

Internal

0	1	2
4	1	5

Summary

- ▶ Various ways to manage Memory.
- ▶ It is always a trade-off between speed and efficiency.
- ▶ It has a **big impact** on the real performance.

Reference:

- ▶ Goodrich and Tamassia, 646–648 and 668–669
- ▶ Aho and Ullman, 146–159
- ▶ most OS books.

Outline

Memory Management

- Principle of Locality

- Block Replacement Strategies

Heap Management

- Fixed Sized Blocks

- Variable Sized Blocks

- Buddy Block

Final Review

Operators

- ▶ Treat the heap as an ADT
- ▶ Two operators:
 - ▶ $p \leftarrow \text{Allocate}(\textit{size})$
 - ▶ $\text{Free}(p)$

Issues

- ▶ Fixed or variable sized blocks?
- ▶ If fixed, small or large sized blocks?
- ▶ User or OS freeing memory blocks?
Explicit vs Implicit freeing of memory.
- ▶ Time or memory?
- ▶ Can blocks reference each other?

Fixed Sized Blocks

- ▶ Treat memory as an array of blocks
- ▶ Some are in use, some are on a free list
- ▶ User can **only** ask for one block at a time
- ▶ Allocation – grab any block from the free list
- ▶ Free – when no longer needed, return to free list

Fixed Sized Blocks

Issues

- ▶ **Internal Fragmentation**

If user does not need all space in block, the extra space is wasted

- ▶ How do we know when block is no longer needed?

- ▶ **Explicit Free**

Simplest to implement

Fixed Sized Blocks

Implicit Freeing

- ▶ **Implicit Free** with reference counts
 - ▶ Update number of references to an allocated block
 - ▶ Islands of garbage may result
- ▶ **Implicit Free** with garbage collection
 - ▶ Periodically clean memory
 - ▶ 2-pass Mark and Sweep algorithm

Variable Sized Blocks

- ▶ Allocate a contiguous chunk of memory
- ▶ All the issues with fixed sized blocks and more
- ▶ **External Fragmentation**
 - ▶ Division of free space leaves no chunk large enough
 - ▶ May require expensive **Memory Compaction**
- ▶ **Internal Fragmentation**
 - ▶ To simplify allocation, OS may give a larger block than was requested
 - ▶ For example, round up to a multiple of a minimum block size

Variable Sized Blocks

Implementation

- ▶ Maintain an ordered list of each chunk of memory
- ▶ Allocate:
 - ▶ Search for a free hole that is “big enough”
 - ▶ Mark node as allocated, updating neighbour
- ▶ Free:
 - ▶ Mark node as free
 - ▶ Merge with any free neighbours

Variable Sized Blocks

Selection Policies

- ▶ Consider the following sequence on a RAM of size 100:

```
A <- Allocate( 40 )
```


```
B <- Allocate( 40 )
```

```
Free( A )
```


```
C <- Allocate( 10 )
```

- ▶ Various strategies for choosing free block

- ▶ **First Fit**



- ▶ **Best Fit**




Variable Sized Blocks

Comments

- ▶ All strategies can outperform the others in certain cases
- ▶ First-fit and Alternating are simplest to implement
- ▶ Worst-fit needs a priority queue
- ▶ Best-fit needs a more complicated priority queue

Buddy Block


- ▶ Trade less external fragmentation for internal fragmentation
- ▶ OS gives block sizes that are powers of two
 - ▶ Suppose heap has size $N = 2^m$
 - ▶ Valid allocation sizes are $2^0, 2^1, \dots, 2^m$ blocks
 - ▶ Maintain $m + 1$ free lists
 - ▶ Initially all but the 2^m list is empty



Buddy Block

Implementation

- ▶ Only allowed to merge free blocks that are “buddies”
- ▶ Similar to extendible hashing
- ▶ Each block of size 2^k starts at some multiple of 2^k
- ▶ Each block of size 2^k has a buddy of the same size
- ▶ The two buddies are within the same block of size 2^{k+1}



Buddy Block


Operators

- ▶ $p \leftarrow \text{Allocate}(n)$
 - ▶ Round n up to the nearest power of two: 2^i
 - ▶ If no free block of size 2^i , split a block of size 2^{i+1}
 - ▶ This split may require another split ...
- ▶ $\text{Free}(p)$
 - ▶ Add p 's block to the free list of size 2^i
 - ▶ If buddy is in the free list, merge and add to free list of size 2^{i+1}
 - ▶ Again, this may trigger another merge ...


Buddy Block

Example


► A ← Allocate(3)




► B ← Allocate(3)



► C ← Allocate(1)




► D ← Allocate(2)




Buddy Block

Example


► Free(B)



► Free(C)



► Free(D)



Buddy Block

Comments

- ▶ $\Theta(1)$ to find buddy
 - ▶ If free list is implemented with an array
- ▶ $\Theta(m)$ worst case for both Free and Allocate
- ▶ Internal fragmentation wastes approximately 30%
- ▶ External fragmentation still possible

Summary

- ▶ Two extremes and one intermediate solution:
Fixed, Variable Sized Blocks, and **Buddy** Blocks.
- ▶ As for memory management, big impact on performance.

Reference: none in GT nor CLRS.

Outline

Memory Management

- Principle of Locality

- Block Replacement Strategies

Heap Management

- Fixed Sized Blocks

- Variable Sized Blocks

- Buddy Block

Final Review

Final Review

The final can/should cover the following concepts:

- ▶ Algorithmic Analysis
- ▶ ADTs
 - ▶ LIFO, FIFO
 - ▶ Graphs
 - ▶ Trees, ordinal and binary cardinal
 - ▶ Priority Queues and Heaps
 - ▶ Unsorted Dictionaries
 - ▶ Sorted Dictionaries
 - ▶ Valued Dictionaries
- ▶ Applications
 - ▶ Relational Databases and SQL
 - ▶ Compression Algorithms
 - ▶ Pattern Matching Algorithms
 - ▶ Heap and Memory Management.

1. Algorithmic Analysis

- ▶ Asymptotics
- ▶ Sorting algorithms
- ▶ Recursivity (and its analysis)
- ▶ Worst and Average Case complexity
- ▶ Randomized algorithms
- ▶ Lower bounds in the comparison model

2. ADTs

- ▶ LIFO, FIFO
- ▶ Graphs
- ▶ Trees, ordinal and binary cardinal
- ▶ Priority Queues and Heaps
- ▶ Unsorted Dictionaries
 - ▶ with Array
 - ▶ with Lists
- ▶ Sorted Dictionaries
 - ▶ Array
 - ▶ Lists
 - ▶ BST
 - ▶ AVL
 - ▶ (2,4) Trees
 - ▶ B Trees
- ▶ Valued Dictionaries
 - ▶ Hash Table

3. Applications

- ▶ Relational Databases and SQL
- ▶ Compression Algorithms
- ▶ Pattern Matching Algorithms
- ▶ Heap and Memory Management.

Advices

- ▶ Make sure that you know how to do the **assignments**, and the **midterm**.
- ▶ Do the assignments from previous years.
- ▶ Try to remember **Why** and **How** we do things.

Good Luck...